

Explanation-aware vision–language architecture

frozen visual backbone, trainable QLoRA adapters, and span-aware supervision

