

Cross-modal generative alignment for trustworthy visual reasoning

fixed-budget comparison of answer-only and explanation-aware supervision

